

# A Resonance Calculus for Tail-Aware Optimization and Control of Discrete-Event Systems

Christopher A. Stringer  
Independent Researcher, Duluth, GA, USA

November 2025

MSC 2020: 49J20, 90C30, 93C65 (primary); 90B22, 68M20 (secondary)

## Abstract

We propose *Resonance Calculus* (RC) and *Resonance Algebra*, a unified framework for modeling and optimizing the performance of large-scale discrete-event systems under coherence, timing, and tail constraints. RC combines (i) coherence-weighted service models that couple classical service curves with context-dependent alignment factors, (ii) tail-aware risk metrics derived from Extreme Value Theory (EVT) applied to rare high-impact events, and (iii) max-plus algebra for synchronization, cycle-time, and bottleneck analysis on networks of interdependent components. These ingredients yield scalar resonance scores and max-plus eigenvalues that can be used as control signals for routing, scheduling, and resource allocation.

We formulate resonance metrics, relate them to service curves and max-plus cycle times, and outline feedback laws that regulate resonance within prescribed bands under explicit SLO and safety constraints. The framework applies to cloud and edge computing, embedded and IoT controllers, fleet and operations systems, and hybrid quantum–classical workflows. We provide illustrative constructions and discuss how RC can be integrated with existing optimization and control stacks. Long-range speculative physics or propulsion concepts are deliberately excluded; we focus on mechanisms that are concretely implementable with current or near-term technology.

## 1 Introduction

Traditional performance engineering emphasizes central tendencies: mean latency, average utilization, aggregate throughput. In contrast, user-perceived performance and reliability are dominated by three factors:

1. *Coherence*: the degree of alignment between what a system is doing and what the current environment, workload, or users actually require;
2. *Tail behavior*: rare but high-impact events (e.g. the slowest 1% or 0.1% of transactions) that dominate perceived latency, SLO violations, and incident risk;
3. *Timing structure and synchronization*: how delays and misalignment propagate and amplify through networks of queues, services, and controllers.

Large-scale distributed systems exhibit *tail amplification*: even low-probability slowdowns at individual components can affect a substantial fraction of end-to-end requests as fan-out increases. At the same time, workloads are increasingly distributed (microservices, serverless, hybrid cloud), multi-scale (edge devices, IoT, embedded controllers), human-in-the-loop (support centers, on-call operations), and hybrid (quantum– classical workflows, heterogeneous accelerators).

Existing solutions often treat capacity, tails, and synchronization as loosely related but separate concerns. Operators may tune autoscalers against average CPU, inspect tail latency dashboards when incidents occur, and apply ad hoc rules to mitigate cascades. There is no canonical scalar quantity that directly measures how “in tune” a system is with its environment and intent.

Resonance Calculus aims to address this gap. It treats a system as a resonant structure: a network of components whose timing, capacity, and context can be measured and adjusted jointly. Rather than separately optimizing utilization, percentiles, and throughput, RC defines a family of resonance metrics that incorporate:

- coherence between behavior and context;
- tail risk in latency and other risk-heavy metrics;
- timing structure and bottlenecks in dependency graphs.

These metrics drive closed-loop control and scheduling in a way that explicitly targets alignment and tail resilience, not just central tendencies.

## 2 Background and Related Work

Resonance Calculus draws on and connects several established fields.

### 2.1 Network Calculus and Service Curves

Network calculus provides a deterministic framework for worst-case performance analysis of communication networks using arrival curves and service curves [1, 2, 3]. A service curve  $\beta(t)$  typically characterizes the minimum amount of service offered over any interval of length  $t$ . Network calculus yields end-to-end bounds on delay, backlog, and buffer occupancy.

These models usually assume fixed or stationary service and do not explicitly incorporate coherence with higher-level context such as workload composition, user intent, or external constraints. Service guarantees exist, but whether they are applied at the “right” times and places is not captured.

### 2.2 Extreme Value Theory and Tail Modeling

Extreme Value Theory (EVT) provides asymptotically justified models for the tails of distributions, including the Generalized Extreme Value (GEV) families and peaks-over-threshold (POT) approaches based on the Generalized Pareto Distribution (GPD) [4, 5, 6]. In a POT setup, one chooses a high threshold  $u$  and models exceedances  $X - u$  with a GPD, from which high quantiles  $Q_{\text{tail}}(q)$  can be derived for probabilities  $q$  close to 1.

EVT and GPDs are widely used in finance, hydrology, and environmental risk, and increasingly appear in performance engineering and reliability analysis. However, tail modeling is rarely integrated into runtime control loops in production systems; tail percentiles are usually treated as passive observability metrics.

### 2.3 Max-Plus Algebra and Discrete-Event Systems

Max-plus (tropical) algebra uses the operations

$$a \oplus b = \max(a, b), \tag{1}$$

$$a \otimes b = a + b, \tag{2}$$

and has been used extensively to study discrete-event systems (DES) with synchronization but no concurrency [7, 8, 9]. Systems that can be modeled as max-plus linear recursions admit analysis via max-plus eigenvalues and maximum-cycle-mean methods, which capture steady-state cycle times and bottlenecks; see also [10, 11, 12].

These tools are well developed at the theoretical level and in specific application domains, but are seldom exposed as first-class telemetry or control signals in conventional distributed systems.

## 2.4 Synchronization and Phase Oscillators

Synchronization phenomena—in physics, biology, and engineering—are often modeled using populations of coupled phase oscillators [13]. The Kuramoto order parameter

$$R_{\text{phase}}(t) = \frac{1}{N} \sum_{i=1}^N e^{j\theta_i(t)} \quad (3)$$

measures phase coherence, where  $\theta_i(t)$  are oscillator phases. This provides a compact indicator of collective timing behavior, from incoherent to phase-locked regimes.

## 2.5 Tail Latency in Large-Scale Services

Large online services are highly sensitive to latency tails [14]. Even rare hiccups can degrade end-to-end latency for many users, especially in fan-out architectures. Techniques such as hedged requests, redundancy, and admission control illustrate that explicit tail-tolerance is essential, not a secondary concern.

Resonance Calculus can be viewed as a framework that internalizes these lessons into a single algebra: coherence-weighted service, tail-aware models, and timing structure are combined into resonance metrics that directly drive control decisions.

# 3 Resonance Calculus: Core Constructs

We outline the main components of Resonance Calculus. Formal definitions match those in an underlying patent specification; here we focus on structure and intuition.

## 3.1 Coherence-Weighted Service Curves

Let  $\beta(t)$  be a baseline service curve describing minimum service over any interval of length  $t$ . RC introduces a coherence factor  $c(t) \in [0, 1]$  measuring alignment between behavior and context over a corresponding time scale. Example coherence indicators include:

- correlation between workload demand and available capacity;
- fraction of work scheduled in historically high-performance windows;
- match between user intent (priority, deadline) and observed scheduling.

The coherence-weighted service curve is defined as

$$\beta_c(t) = c(t) \beta(t). \quad (4)$$

When  $c(t) \approx 1$ , baseline service is effectively realized; when  $c(t)$  is small, effective service collapses despite nominal capacity. This couples capacity to alignment, not just to resource allocation.

A coherence-weighted service curve (CWSC) module can produce:

- a time series of  $c(t)$  over sliding windows;
- derived latency or backlog bounds based on  $\beta_c(t)$ ;
- scalar coherence scores summarizing alignment over given horizons.

### 3.2 Tail-Event Modeling via EVT

For a performance metric  $X$  (e.g. response time, queueing delay, cost or energy spikes), RC collects samples and selects a high threshold  $u$ . For exceedances  $X > u$ , it considers  $Y = X - u$  and fits a GPD:

$$\Pr(Y > y) = \left(1 + \xi \frac{y}{\sigma}\right)^{-1/\xi}, \quad y > 0, \quad (5)$$

with shape  $\xi$  and scale  $\sigma$ . Tail quantiles are then

$$Q_{\text{tail}}(q) = u + \frac{\sigma}{\xi} \left( (1 - q)^{-\xi} - 1 \right), \quad q \in (0, 1), \quad (6)$$

for high  $q$  (e.g. 0.99, 0.999).

RC converts tail quantiles and parameters into tail-health scores by normalizing against SLOs or reference levels and mapping to a scale where lower delays correspond to higher scores. These scores influence resonance metrics and guardrail logic.

### 3.3 Resonance Algebra and Max-Plus Timing

RC represents system components as nodes in a directed graph:

- nodes may denote services, queues, processes, devices, robots, or other agents;
- edges denote dependencies or routes with baseline delays  $d_{ij}$ ;
- each node  $i$  has a coherence factor  $c_i$ .

For each edge  $(i, j)$ , RC defines a coherence-modulated delay:

$$\tilde{d}_{ij} = d_{ij} - \gamma \min(c_i, c_j), \quad \gamma \geq 0. \quad (7)$$

The collection of  $\tilde{d}_{ij}$  values forms a max-plus matrix  $\tilde{D}$ , where addition is replaced by  $\oplus = \max$  and multiplication by  $\otimes = +$ .

A max-plus eigenvalue  $\lambda_{\text{res}}$  satisfies

$$\tilde{D} \otimes x = \lambda_{\text{res}} \otimes x \quad (8)$$

for some eigenvector  $x$ . In practice,  $\lambda_{\text{res}}$  can be computed via maximum-cycle-mean algorithms:

$$\lambda_{\text{res}} = \max_C \frac{w(C)}{|C|}, \quad (9)$$

where  $C$  ranges over cycles,  $w(C)$  is the sum of weights along  $C$ , and  $|C|$  is its length.

$\lambda_{\text{res}}$  acts as a resonance-sensitive cycle time or inverse throughput indicator: it encapsulates timing structure and bottlenecks under coherence-adjusted conditions.

## 4 Resonance Scores and Control

### 4.1 Subsystem and Global Resonance

For a subsystem  $i$ , RC defines:

- a system score  $S_i(t)$ , derived from coherence-weighted service, tail-health, and timing metrics;
- a human/operator score  $H_i(t)$  (e.g. satisfaction, overload indicators);
- a context score  $C_i(t)$  (e.g. environmental or external conditions).

Subsystem resonance is

$$R_i(t) = w_{S,i}S_i(t) + w_{H,i}H_i(t) + w_{C,i}C_i(t), \quad w_{S,i}, w_{H,i}, w_{C,i} \geq 0. \quad (10)$$

Global resonance aggregates subsystems:

$$R_{\text{global}}(t) = \frac{\sum_i w_i R_i(t)}{\sum_i w_i}, \quad (11)$$

where  $w_i$  captures importance or scale. These scalars provide compact, tunable measures of system “harmony” at multiple levels.

### 4.2 Learning Resonance Weights

Given features  $f_k(t)$  (e.g. utilization, tail scores, coherence indicators, human feedback) and a target trajectory  $R^*(t)$  (e.g. labeled “good states”), RC can learn non-negative weights  $w_k$  via:

$$\min_{w \geq 0} \sum_t \left( R^*(t) - \sum_k w_k f_k(t) \right)^2 + \lambda \|w\|_1, \quad (12)$$

where  $\lambda$  controls sparsity. This yields an interpretable resonance model that aligns with observed success.

### 4.3 Feedback and Control

At a coarse time scale, resonance dynamics can be approximated by

$$\frac{dR}{dt} = \alpha(C - R) - \beta E, \quad (13)$$

where  $R$  is a resonance metric (e.g.  $R_{\text{global}}$ ),  $C$  is an input coherence signal (structural improvements, favorable conditions),  $E$  is a disturbance (shocks, spikes, failures), and  $\alpha, \beta > 0$ .

Controllers can use familiar proportional–integral structures:

$$u(t) = k_p(R^* - R(t)) + k_i \int_0^t (R^* - R(\tau)) d\tau, \quad (14)$$

where  $u(t)$  is a control signal and  $k_p, k_i \geq 0$  are gains. Control actions include adjusting:

- scheduling intervals and priority weights;
- resource allocations (CPU, memory, bandwidth);
- request-routing rules and replica counts;
- batch sizes and micro-delays for deferrable workloads.

Guardrails enforce SLOs and safety constraints and ensure stable behavior.

## 5 Synchronization Intelligence and Phase Metrics

### 5.1 Phase Oscillators and Global Coherence

RC can model each computational target as a phase oscillator with phase  $\theta_i(t)$  derived from timing signals (arrivals, completions, timers). The global phase coherence is

$$R_{\text{phase}}(t) = \frac{1}{N} \sum_{i=1}^N e^{j\theta_i(t)}, \quad (15)$$

with  $|R_{\text{phase}}(t)| \approx 0$  indicating desynchronization and  $|R_{\text{phase}}(t)| \approx 1$  strong phase-locking.

Phase dispersion

$$\sigma_\theta(t) = 1 - |R_{\text{phase}}(t)| \quad (16)$$

becomes a control signal.

### 5.2 Spectral Observability

A spectral observer applies time-frequency analysis (e.g. short-time Fourier transform, wavelets) to event time series (queue lengths, arrivals, completions) to compute:

- dominant frequencies (periodicities in workload or behavior);
- spectral entropy  $H_s(t)$ , indicating rigidity or noise in timing;
- phase-locking values between components.

Low entropy can indicate overly rigid behavior; high entropy may signal unstructured timing.

### 5.3 Adaptive Coupling and Resonance Bands

An adaptive coupling module adjusts a coupling coefficient  $K(t)$ , for example:

$$K(t) = \text{clip}(K_0 + a \sigma_\theta(t) - b \text{risk}_{p99}(t), K_{\min}, K_{\max}), \quad (17)$$

where  $\sigma_\theta(t)$  measures phase dispersion,  $\text{risk}_{p99}(t)$  is a tail-risk indicator, and  $K_0, a, b, K_{\min}, K_{\max}$  are parameters.

A resonance-band controller maintains a resonance metric  $R(t)$  within  $[R_{\text{low}}, R_{\text{high}}]$  by:

- increasing coupling or applying coordinated nudges when  $R(t)$  is too low;
- injecting small randomized timing variations or reducing coupling when  $R(t)$  is too high and entropy is low.

This prevents both chaotic desynchronization and brittle over-synchronization.

## 6 SyncScript: Human-Facing Resonance

Many systems are human-in-the-loop: operators, teams, knowledge workers, and on-call engineers. RC exposes a human-facing layer, SyncScript, that lets humans and applications express work in a resonance-aware format.

## 6.1 Task Model

SyncScript represents tasks with fields such as:

- title, description;
- type (e.g. deep work, shallow work, admin, meeting, creative);
- duration estimate;
- importance (e.g. low, normal, high, critical);
- deadline, maximum deferral;
- batchable, maximum batch latency;
- phase affinity (e.g. morning, early afternoon, low-resonance windows).

The resonance engine uses this metadata and its coherence landscape over time to place:

- deep work in high-resonance windows;
- administrative tasks in low-resonance valleys;
- overdue or critical items where they alleviate tail risk without destabilizing the system.

## 6.2 Illustrative DSL Snippet

A simple SyncScript-like notation might be:

```
task "Deep Work { Resonance Paper" {
    type: deep_work
    duration: 90m
    importance: high
    deadline: 2025-11-30
    phase_affinity: [morning, early_afternoon]
    batchable: false
}

task "Admin Inbox Sweep" {
    type: shallow_work
    duration: 30m
    importance: normal
    batchable: true
    max_batch_latency: 2d
    phase_affinity: [low_resonance]
}

routine "Daily SyncScript Loop" {
    triggers: [start_of_day]
    include: [all_overdue, all_today, high_importance]
}
```

SyncScript provides a concrete bridge between abstract resonance metrics and day-to-day decision making.

## 7 Example Use Cases

### 7.1 Cloud Microservice Architectures

Services or containers are nodes in the max-plus graph; edges encode RPC calls or data flows. Coherence reflects alignment between request mix, capacity, and SLOs. EVT models tail latencies;  $\lambda_{\text{res}}$  reflects bottlenecks.

RC can:

- adjust autoscaling policies based on resonance metrics rather than raw CPU;
- modify routing rules when tail-health degrades on specific paths;
- tune batch sizes and micro-delays to reduce contention bursts.

### 7.2 Edge and IoT Control Systems

Edge devices operate under constraints of energy, connectivity, and duty cycles. RC can:

- compute coherence between sensing, communication, and actuation schedules;
- model tail behavior in communication and processing delays;
- adjust sampling frequency, transmit windows, and local processing strategies to keep events aligned with constraints while mitigating worst-case delays.

### 7.3 Fleet and Operations Management

Vehicles or robots are modeled as nodes; routes define edges. Coherence measures how positions and assignments match spatial and temporal demand. Tail metrics capture worst-case delivery or response times.

RC can recommend:

- rebalancing vehicles;
- pre-positioning resources ahead of predicted demand;
- adjusting service territories to avoid persistent low-resonance zones.

### 7.4 Hybrid Quantum–Classical Orchestration

Hybrid workloads involve classical preprocessing, quantum kernels, and classical post-processing. RC:

- represents quantum and classical stages as nodes in the graph;
- uses coherence factors to capture alignment between job types, hardware calibration windows, and coherence-time forecasts;
- applies EVT to model tail latencies or error bursts;
- schedules critical quantum operations into predicted high-coherence windows and shapes non-critical classical work around them.

Here RC acts as an orchestration layer that respects both classical constraints (queues, CPUs, networks) and quantum-specific constraints (coherence times, calibration schedules), without making speculative physics claims.

## 8 Observability, Guardrails, and Benchmarking

### 8.1 Telemetry and Explainability

RC exports:

- resonance metrics (global and per-subsystem);
- tail metrics (EVT parameters, tail quantiles, tail-health scores);
- spectral metrics (dominant frequencies, spectral entropy, phase coherence);
- controller actions and state.

An explanation interface can, for a given decision, report:

- whether a task was executed, deferred, or batched;
- contributing features and thresholds;
- current resonance band and guardrail state.

### 8.2 SLOs and Safety Constraints

RC operates under explicit SLOs and safety conditions, such as:

- latency SLOs (e.g.  $p_{99}$  bounds for key routes);
- bounds on acceptable changes in throughput or error rates;
- limits on deferral and batching for specific classes of work.

Guardrails can:

- disable resonance-based control and revert to baseline when performance degrades beyond thresholds compared to control groups;
- reduce control strength when oscillatory behavior or instability is detected;
- trigger alerts when resonance metrics or tail indicators enter pathological regimes.

### 8.3 Benchmarking Harness

A benchmarking harness can:

- generate controlled workloads with bursts, diurnal patterns, adversarial patterns;
- run the system in baseline and resonance-enabled modes;
- collect latency distributions, throughput, error rates, resource overhead, and resonance trajectories;
- produce comparative reports quantifying improvements in tail performance, coherence, and synchronization.

This provides empirical evidence for the benefits and limits of RC in concrete environments.

## 9 Conclusion

Resonance Calculus and Resonance Algebra provide a unified framework for coherence-aware, tail-aware, and timing-aware optimization in complex systems. By coupling coherence-weighted service models, EVT-based tail modeling, and max-plus timing analysis with phase and spectral observability, RC yields resonance metrics that can directly drive control and scheduling decisions.

This paper has focused deliberately on mechanisms that are concretely implementable in current or near-term systems: distributed computing, edge and IoT control, fleet and operations management, and hybrid quantum-classical orchestration. More speculative directions in physics, propulsion, and long-horizon technologies are intentionally excluded and reserved for future work.

Near-term efforts will concentrate on:

- reference implementations and SDKs;
- integration with existing observability and control stacks;
- benchmarking RC against baseline control across realistic workloads;
- developing SyncScript and related human-facing tools for resonance-aware workflows.

In this sense, Resonance Calculus is proposed not as a speculative physical theory, but as a practical engineering framework: a way to make systems and organizations move in rhythm with the environments they inhabit.

## References

- [1] R. L. Cruz. A calculus for network delay, Part I: Network elements in isolation. *IEEE Transactions on Information Theory*, 37(1), 1991.
- [2] J.-Y. Le Boudec. *Network Calculus*. Springer, 2001; and related tutorial material, 2019.
- [3] A. Van Bemten and W. Kellerer. Network Calculus: A Comprehensive Guide. Technical Report 201603, Technical University of Munich, 2016.
- [4] S. Coles. *An Introduction to Statistical Modelling of Extreme Values*. Springer Series in Statistics, 2001.
- [5] P. Embrechts, C. Klüppelberg, and T. Mikosch. *Modelling Extremal Events for Insurance and Finance*. Springer, 1997.
- [6] M. Falk. Peaks-over-threshold stability of multivariate generalized Pareto distributions. *Journal of Multivariate Analysis*, 2008.
- [7] J. Komenda et al. Max-plus algebra and discrete event systems, 2017.
- [8] B. De Schutter. Analysis and control of max-plus linear discrete-event systems: an introduction. *Annual Reviews in Control*, 2020.
- [9] G. Cohen et al. Max-plus algebra and system theory: where we are and where to go now. *Annual Reviews in Control*, 1999.
- [10] L. Hardouin. Control and state estimation for max-plus linear systems. *Foundations and Trends in Systems and Control*, 2018.

- [11] A. Gupta. A framework for studying stability of switching max-plus linear systems. *arXiv preprint*, arXiv:2007.02818, 2020.
- [12] N. K. Krivulin. Max-plus algebra models of queueing networks. *arXiv preprint*, arXiv:1212.0578, 2012.
- [13] J. A. Acebrón et al. The Kuramoto model: A simple paradigm for synchronization phenomena. *Reviews of Modern Physics*, 77, 2005.
- [14] J. Dean and L. A. Barroso. The Tail at Scale. *Communications of the ACM*, 56(2), 2013.